

Kubernetes Notes – (Namespaces to HPA/VPA)

Table of Contents

1. [Namespaces](#)
 2. [Pods](#)
 3. [ReplicaSets](#)
 4. [Deployments](#)
 5. [Services \(ClusterIP, NodePort, LoadBalancer\)](#)
 6. [Resource Requests and Limits](#)
 7. [Horizontal Pod Autoscaler \(HPA\)](#)
 8. [Vertical Pod Autoscaler \(VPA\)](#)
 9. [Key Differences between HPA and VPA](#)
-

Namespaces

Namespaces in Kubernetes are like virtual clusters. They allow you to divide cluster resources between multiple users (via resource quota).

Example

```
apiVersion: v1
```

```
kind: Namespace
```

```
metadata:
```

```
  name: dev
```

```
kubectl create -f namespace.yaml
```

```
kubectl get namespaces
```

```
kubectl config set-context --current --namespace=dev
```

Pods

A Pod is the smallest deployable unit in Kubernetes. It can contain one or more containers that share storage/network.

Example

```
apiVersion: v1
kind: Pod
metadata:
  name: nginx-pod
  labels:
    app: nginx
spec:
  containers:
    - name: nginx
      image: nginx:latest
      ports:
        - containerPort: 80
```

kubectl apply -f pod.yaml

kubectl get pods

ReplicaSets

ReplicaSet ensures a specified number of pod replicas are running at all times.

Example

```
apiVersion: apps/v1
kind: ReplicaSet
metadata:
  name: nginx-replicaset
spec:
  replicas: 3
  selector:
    matchLabels:
      app: nginx
  template:
    metadata:
      labels:
```

app: nginx

spec:

containers:

- name: nginx

image: nginx

Deployments

Deployments manage ReplicaSets and provide declarative updates for Pods and ReplicaSets.

Example

apiVersion: apps/v1

kind: Deployment

metadata:

name: nginx-deployment

spec:

replicas: 2

selector:

matchLabels:

app: nginx

template:

metadata:

labels:

app: nginx

spec:

containers:

- name: nginx

image: nginx

kubectl apply -f deployment.yaml

kubectl rollout status deployment nginx-deployment

Services

Services expose a set of pods to other services inside/outside the cluster.

◆ Types of Services

1. **ClusterIP** – default; accessible within the cluster.
 2. **NodePort** – accessible via <NodeIP>:<NodePort>.
 3. **LoadBalancer** – external IP for cloud.
-

✅ Example: NodePort

apiVersion: v1

kind: Service

metadata:

name: nginx-service

spec:

type: NodePort

selector:

app: nginx

ports:

- port: 80

targetPort: 80

nodePort: 30080

⚙️ Resource Requests and Limits

Kubernetes allows setting CPU and Memory resources for containers.

- **Request** = minimum guaranteed
- **Limit** = maximum allowed

✅ Example

resources:

requests:

memory: "64Mi"

cpu: "250m"

limits:

memory: "128Mi"

cpu: "500m"

Horizontal Pod Autoscaler (HPA)

HPA automatically scales the number of pods in a deployment based on observed CPU/memory utilization.

Example

apiVersion: autoscaling/v2

kind: HorizontalPodAutoscaler

metadata:

name: nginx-hpa

spec:

scaleTargetRef:

apiVersion: apps/v1

kind: Deployment

name: nginx-deployment

minReplicas: 1

maxReplicas: 5

metrics:

- type: Resource

resource:

name: cpu

target:

type: Utilization

averageUtilization: 50

kubectl autoscale deployment nginx-deployment --cpu-percent=50 --min=1 --max=5

Vertical Pod Autoscaler (VPA)

VPA automatically adjusts CPU and memory **requests** and **limits** for containers.

✓ Example

apiVersion: autoscaling.k8s.io/v1

kind: VerticalPodAutoscaler

metadata:

name: nginx-vpa

spec:

targetRef:

apiVersion: "apps/v1"

kind: Deployment

name: nginx-deployment

updatePolicy:

updateMode: "Auto"

Note: VPA may restart pods to apply new resource values.

🔄 HPA vs VPA

Feature	HPA	VPA
Scales	Number of pods	CPU/memory of pods
Based On	Metrics like CPU/Memory Historical + live resource usage	
Impact	More pods	Pod restarts
Use Case	Stateless apps	Stateful apps (like DB)
Update Mode	Continuous	Off/Auto/Initial

✓ Summary Commands

Create all from YAMLs

kubectl apply -f <filename>.yaml

Check HPA

kubectl get hpa

Check VPA (if vpa components installed)

```
kubectl get vpa
```

```
# Set context to namespace
```

```
kubectl config set-context --current --namespace=dev
```



Pro Tip: Use `kubectl describe <resource>` to debug deeper!