

What is Kubernetes (K8s)?

Kubernetes is a container orchestration platform. It automates the deployment, scaling, and management of containerized applications across a cluster of machines.

High-Level Architecture

Kubernetes has two main layers:

1. **Control Plane (Master Node)** – manages the cluster
 2. **Worker Nodes** – runs application workloads
-

1. Control Plane Components (Master Node)

The control plane makes global decisions about the cluster (e.g., scheduling), and detects and responds to cluster events (e.g., restarting failed pods).

Components:

a. kube-apiserver

- The front-end of the Kubernetes control plane.
- Accepts REST calls from CLI (kubectl) or UI.
- Authenticates and validates requests.
- Communicates with etcd.

b. etcd

- A distributed key-value store.
- Stores all cluster data (e.g., nodes, pods, configs).
- Backup of etcd = backup of your whole cluster state.

c. kube-scheduler

- Watches for new pods with no assigned node.
- Selects the best node to run the pod based on:
 - Resource availability
 - Node affinity
 - Taints and tolerations
 - Other constraints

d. kube-controller-manager

- Runs various controllers:
 - **Node Controller** – monitors node status
 - **Replication Controller** – ensures desired pod count
 - **Endpoints Controller** – manages endpoint objects
 - **Service Account & Token Controller**

e. cloud-controller-manager (*optional*)

- Integrates with cloud provider APIs.
 - Manages:
 - Load balancers
 - Volumes
 - Node instances
-

2. Worker Node Components

Worker nodes run the containers that make up your application.

Components:

a. kubelet

- Communicates with API Server.
- Ensures the containers described in PodSpecs are running and healthy.

b. kube-proxy

- Manages network rules.
- Enables communication between services (internal & external).
- Uses iptables or IPVS.

c. Container Runtime

- The software responsible for running containers.
 - Examples: **Docker, containerd, CRI-O**
-

3. Pods and Workloads

What is a Pod?

- The smallest and simplest unit in Kubernetes.

- Encapsulates one or more containers with shared storage/network.
 - Containers in the same pod can communicate via localhost.
-

Full Working Flow (End-to-End)

1. **User runs a command:** `kubectl apply -f app.yaml`
 2. **API Server receives the request**
 3. **Validates and stores it in etcd**
 4. **Scheduler assigns a pod to a node**
 5. **Controller manager ensures desired state**
 6. **kubelet on selected node creates the pod**
 7. **Container runtime pulls and starts the container**
 8. **kube-proxy sets up networking for service discovery**
-

Networking Model in K8s

- Each Pod gets a **unique IP address**.
- Pods can talk to each other **without NAT**.
- Services (ClusterIP, NodePort, LoadBalancer) expose Pods.
- kube-dns or CoreDNS allows name-based resolution.